

The Art of Neural Nets

Marco Tavora

marcotav65@gmail.com



Preamble

The challenge of recognizing artists given their paintings has been, for a long time, far beyond the capability of algorithms.

Recent advances in deep learning, specifically the development of convolutional neural networks, have made that task possible.



Challenge

Given a painting can machines identify the painter?



Can you identify these
painter or maybe their
period/style?



Challenge

Given a painting can machines identify the painter?



Woman in Hat [Olga] by
Pablo Picasso, 1923

Portrait of Olga Picasso, by
Pablo Picasso, 1935

Both are from Picasso
(1923, 1935)




Challenge

Complex but useful problem:

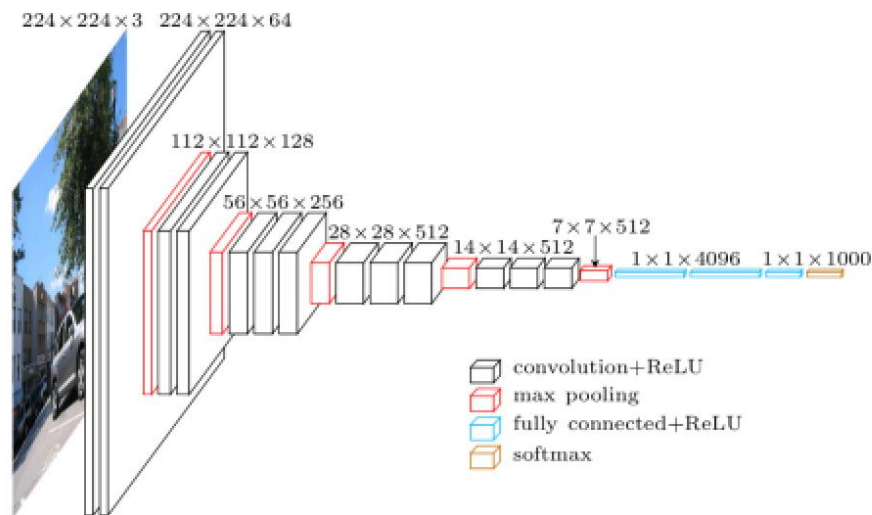
- Painters change styles along their career
- It is not just object or face recognition (painters can paint many different things)
- Methods employed by specialists can be invasive and cause damage do the painting
- High performance classifiers can also help identify forgeries

Deep learning to the rescue!



Deep learning

- DL methods of image processing can help.
- Most previous work on artist identification start by choosing the painters' characteristics as features
- Since NN do not need “feature hand-crafting” they are ideal for this task

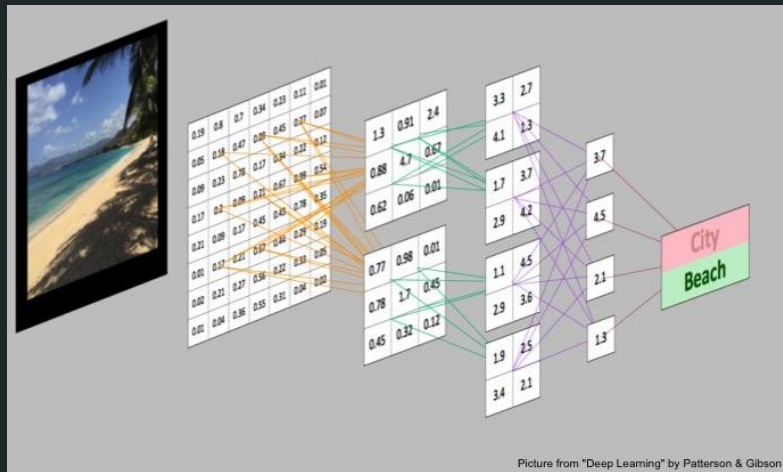


VGGNet Architecture

Methodology: Convolutional Neural Networks (CNNs)

Particularly proficient at image classification

They make assumption that inputs are images and that vastly reduces number of parameters






Picture from "Deep Learning" by Patterson & Gibson


Methodology: Convolutional Neural Networks (CNNs)

Layers are feature identifiers:

Convolutional layers filter input and they identify edges, curves, ...




Pooling layers shrink input dimension by an integer factor



Procedure

- Multi-class classification problem
- Painters must have a minimum number of paintings in the training set to ensure sufficient sample size
- Running time increases when the number of classes increases, so I ended up choosing 3 painters (contrast with 1 painter = binary classification)

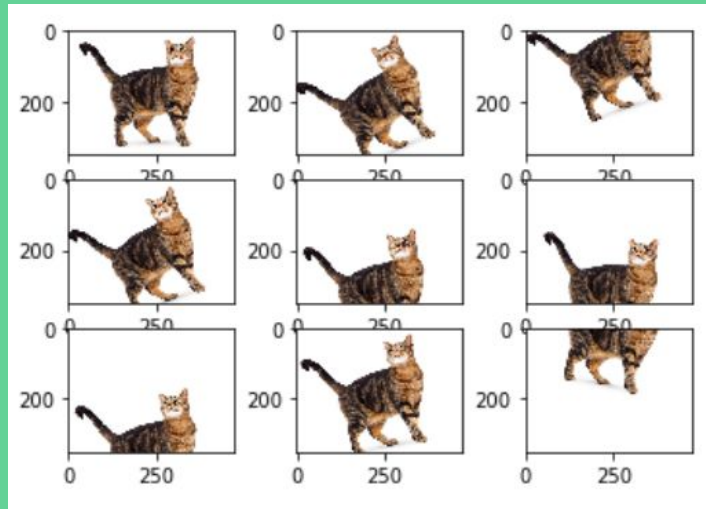


Procedure

- We have a small sample (for image recognition standards)
- Tactics to attack this problem of lack of data:
 - Image augmentation
 - **Transfer learning**

Image Augmentation

- Combats high expense of getting new data by creating new data from existing data
- Keras very easily manipulates images to get many altered versions of the same image
- Modifications include rotations, shifts, shears, flips, whitening and others




What is Transfer Learning?

- When knowledge gained during training in one type of problem is used to train in other type of problem (e.g. knowledge gained learning to recognize cars applies when trying to recognize trucks)
- TL and CNNs: nowadays, training CNNs from scratch is quite rare because of sample size limitations
- Pre-train a CNN on a large dataset (e.g. ImageNet with 1.2M images, 1K categories) and apply it to our task


Feature Extractor

- We remove only the last full-connected layer of the pre-trained CNN and treat the rest as a fixed feature extractor
- I replace the final fully-connected layer with a new one to obtain a score for each painter in my dataset instead of a score for ImageNet classes.



Feature Extractor

To verify if features from Google's InceptionV3 were useful as a starting point, I used it directly on my training data



Inception V3 in Keras

```
applications.InceptionV3(include_top=False, weights='imagenet').summary()
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_2 (InputLayer)	(None, None, None, 3 0		
conv2d_95 (Conv2D)	(None, None, None, 3 864		input_2[0][0]
batch_normalization_95 (BatchNo	(None, None, None, 3 96		conv2d_95[0][0]
activation_95 (Activation)	(None, None, None, 3 0		batch_normalization_95[0][0]
conv2d_96 (Conv2D)	(None, None, None, 3 9216		activation_95[0][0]
batch_normalization_96 (BatchNo	(None, None, None, 3 96		conv2d_96[0][0]
activation_96 (Activation)	(None, None, None, 3 0		batch_normalization_96[0][0]
conv2d_97 (Conv2D)	(None, None, None, 6 18432		activation_96[0][0]
batch_normalization_97 (BatchNo	(None, None, None, 6 192		conv2d_97[0][0]
activation_97 (Activation)	(None, None, None, 6 0		batch_normalization_97[0][0]
max_pooling2d_5 (MaxPooling2D)	(None, None, None, 6 0		activation_97[0][0]
conv2d_98 (Conv2D)	(None, None, None, 8 5120		max_pooling2d_5[0][0]
batch_normalization_98 (BatchNo	(None, None, None, 8 240		conv2d_98[0][0]
activation_98 (Activation)	(None, None, None, 8 0		batch_normalization_98[0][0]

Feature Extractor



```
preds = base_model.predict(x)
print('Predicted:', decode_predictions(preds))
```


Predicted: [[('n02980441', 'castle', 0.14958656), ('n03788195', 'mosque', 0.10087459), ('n03445777', 'church', 0.086444855), ('n03388043', 'fountain', 0.07997718), ('n09428293', 'seashore', 0.07918877)]]

Implementation

Metrics: I looked at the accuracy and (binary cross-entropy) loss for both training and validation set over the training epochs

I only used rectified linear activation functions (ReLU) and the Adam optimizers

$$J = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$



Implementation

My fully-connected top model had convolutional layers with dimensions 512, 256, 128,..., 16 and 8 with 50% drop-outs in between. Adding layers was very helpful with both accuracy and losses.


Then I inserted a small CNN on top of the pre-trained network and trained it on those features get the classification I needed.

```
model.summary()
```

Layer (type)	Output Shape	Param #
flatten_6 (Flatten)	(None, 8192)	0
dense_41 (Dense)	(None, 512)	4194816
dropout_36 (Dropout)	(None, 512)	0
dense_42 (Dense)	(None, 256)	131328
dropout_37 (Dropout)	(None, 256)	0
dense_43 (Dense)	(None, 128)	32896
dropout_38 (Dropout)	(None, 128)	0
dense_44 (Dense)	(None, 64)	8256
dropout_39 (Dropout)	(None, 64)	0
dense_45 (Dense)	(None, 32)	2080
dropout_40 (Dropout)	(None, 32)	0
dense_46 (Dense)	(None, 16)	528
dropout_41 (Dropout)	(None, 16)	0
dense_47 (Dense)	(None, 8)	136
dropout_42 (Dropout)	(None, 8)	0
dense_48 (Dense)	(None, 3)	27
Total params: 4,370,067		
Trainable params: 4,370,067		
Non-trainable params: 0		

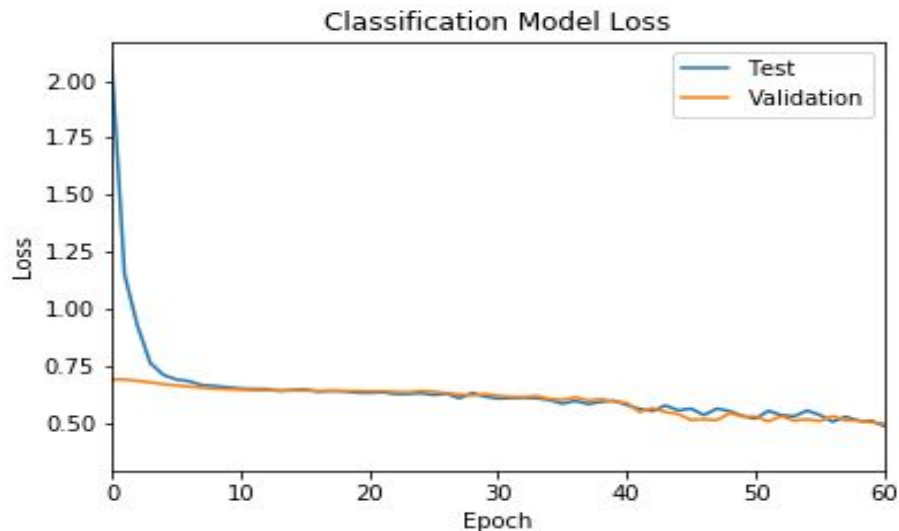
Metrics I

Accuracy for each epoch approaches 76% and baseline is 33.3%. Also train and validation accuracies track very close indicating little overfitting



Metrics II

Loss for each epoch - training and validation are also very close together (binary cross-entropy)



Conclusion I

- I used CNNs to identify the artist of a given painting. The dataset contains (after filtering) around 2000 paintings by 3 painters. Code structure allows this number to be any value (I started with 27 painters but it was too slow)
- I used transfer learning, fine-tuning the top layers of a deep pre-trained networks (Inception V3) and obtained:
 - reasonably high accuracy 76%
 - No overfitting

Conclusion II

- Confusion matrix analysis (partially done but not included in the slides)
- More thorough analysis using the `KerasClassifier()` wrapper for a classification neural network created in Keras and used in the scikit-learn library (`keras.wrappers.scikit_learn`)
- More painters, more paintings, larger networks, etc